



deeplearning.ai


# Convolutional Neural Networks

---

## Computer vision

# Computer Vision Problems

## Image Classification



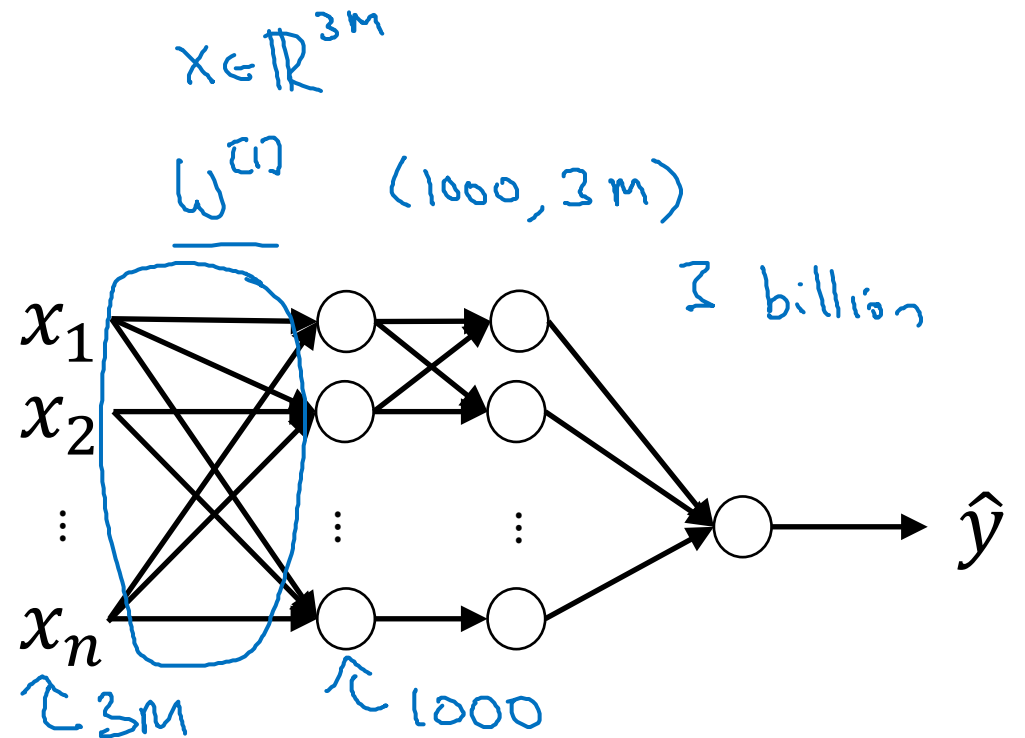

## Neural Style Transfer



## Object detection



# Deep Learning on large images







deeplearning.ai

# Convolutional Neural Networks

---

Edge detection  
example

# Computer Vision Problem




# Vertical edge detection

$$\rightarrow 3 \times 1 + 1 \times 1 + 2 \times 1 + 0 \times 0 + 5 \times 0 + 7 \times 0 + 1 \times -1 + 8 \times -1 + 2 \times -1 = -5$$

3 <sup>1</sup>	0 <sup>1</sup>	1 <sup>-1</sup>	2 <sup>-1</sup>	7 <sup>-0</sup>	4 <sup>-1</sup>
1 <sup>1</sup>	5 <sup>1</sup>	8 <sup>-1</sup>	9 <sup>-1</sup>	3 <sup>-0</sup>	1 <sup>-1</sup>
2 <sup>1</sup>	7 <sup>1</sup>	2 <sup>-1</sup>	5 <sup>-1</sup>	1 <sup>-0</sup>	3 <sup>-1</sup>
0 <sup>1</sup>	1 <sup>1</sup>	3 <sup>-1</sup>	1 <sup>-1</sup>	7 <sup>-0</sup>	8 <sup>-1</sup>
4	2	1	6	2	8
2	4	5	2	3	9

6x6



=


-5	-4	0	8
-10	-2	2	3
0	-2	-4	-7
-3	-2	-3	-16

4x4

# Vertical edge detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0

6x6




\*

1	0	-1
1	0	-1
1	0	-1

3x3


\*



=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0

4x4







deeplearning.ai

# Convolutional Neural Networks

---

More edge  
detection



# Vertical edge detection examples

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0



\*

1	0	-1
1	0	-1
1	0	-1



=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10



\*

1	0	-1
1	0	-1
1	0	-1




=

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0




# Vertical and Horizontal Edge Detection



1	0	-1
1	0	-1
1	0	-1

Vertical



1	1	1
0	0	0
-1	-1	-1

Horizontal

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

6x6

\*



1	1	1
0	0	0
-1	-1	-1

=

0	0	0	0
30	10	-10	-30
30	10	-10	-30
0	0	0	0

# Learning to detect edges

1	0	-1
1	0	-1
1	0	-1



1	0	-1
2	0	-2
1	0	-1



Sobel filter

3	0	-3
10	0	-10
3	0	-3

Scharr filter



3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

convolution  
\*

$W_1$	$W_2$	$W_3$
$W_4$	$W_5$	$W_6$
$W_7$	$W_8$	$W_9$

3x3

=

45°  
70°  
73°






deeplearning.ai

# Convolutional Neural Networks

---



## Padding

# Padding


- Shrinky output
- throw away info from edge

Padding prior to convolution helps:

- avoiding shrinking activations
- exploiting more the info in the edges of the image (because if we pad, the edge pixels will be used in more many convolution operations).



=



$$\frac{6 \times 6}{n \times n} \rightarrow 8 \times 8$$

$$p = \text{padding} = \underline{1}$$

$$n - f + 1 \times n - f + 1$$

$$6 - 3 + 1 = 4$$

$$n + 2p - f + 1 \times n + 2p - f + 1$$

$$6 + 2 - 3 + 1 \times \underline{\underline{6}} = 6 \times 6$$

# Valid and Same convolutions

→ no padding

“Valid”:  $n \times n \quad * \quad f \times f \quad \rightarrow \quad \frac{n-f+1}{1} \times n-f+1$

$6 \times 6 \quad * \quad 3 \times 3 \quad \rightarrow \quad 4 \times 4$

“Same”: Pad so that output size is the same as the input size.

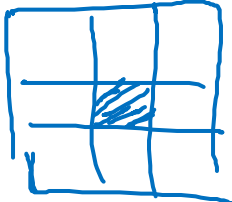
$$n+2p-f+1 \times n+2p-f+1$$
$$\cancel{n+2p-f+1} = \cancel{n} \Rightarrow \boxed{p = \frac{f-1}{2}}$$

$3 \times 3 \quad p = \frac{3-1}{2} = 1$

$5 \times 5 \quad f=5$

$f$  is usually odd

$1 \times 1$   
 $3 \times 3$   
 $5 \times 5$   
 $7 \times 7$



$p=2$



deeplearning.ai

# Convolutional Neural Networks

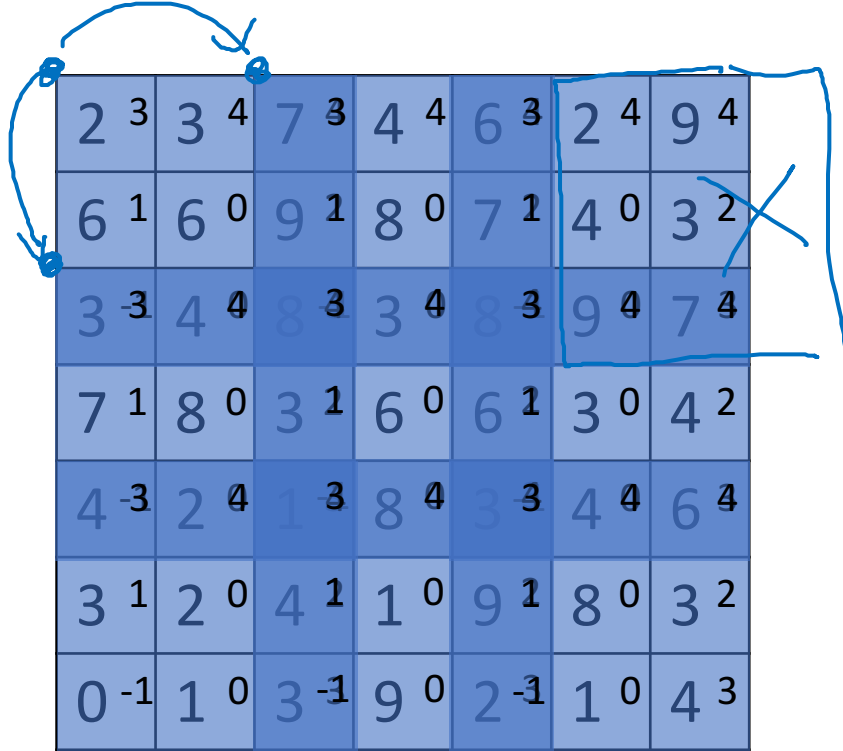
---

## Strided convolutions



# Strided convolution

- to reduce activation size
- to reduce computational load



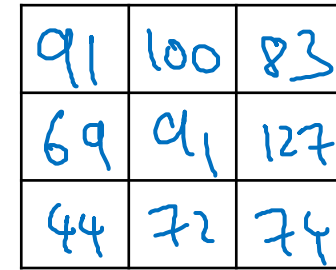
2	3	3	4	7	3	4	4	6	3	2	4	9	4
6	1	6	0	9	1	8	0	7	1	4	0	3	2
3	3	4	4	8	3	3	4	8	3	9	4	7	4
7	1	8	0	3	1	6	0	6	1	3	0	4	2
4	3	2	4	1	3	8	4	3	3	4	4	6	4
3	1	2	0	4	1	1	0	9	1	8	0	3	2
0	-1	1	0	3	-1	9	0	2	-1	1	0	4	3

\*

3	4	4
1	0	2
-1	0	3

3x3

=



91	100	83
69	91	127
44	72	74

3x3

Stride = 2

$\lfloor z \rfloor = \text{floor}(z)$

$n \times n$  \*  $f \times f$   
 padding  $p$       stride  $s$   
 $s = 2$

$$\left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor$$

$$\frac{7 + 0 - 3}{2} + 1 = \frac{4}{2} + 1 = 3$$

# Summary of convolutions

$n \times n$  image       $f \times f$  filter

padding  $p$       stride  $s$

Output Size:

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \underbrace{\frac{n+2p-f}{s}} + 1 \right\rfloor$$

# Technical note on cross-correlation vs. convolution


## Convolution in math textbook:

2 <sup>7</sup>	3 <sup>2</sup>	7 <sup>5</sup>	4	6	2
6 <sup>9</sup>	6 <sup>0</sup>	9 <sup>4</sup>	8	7	4
3 <sup>-1</sup>	4 <sup>1</sup>	8 <sup>3</sup>	3	8	9
7	8	3	6	6	3
4	2	1	8	3	4
3	2	4	1	9	8

3	4	5
1	0	2
-1	9	7

7	2	5
9	0	4
-1	1	3

- In mathematics, the convolution operation requires preliminarily flipping the filter. In ML this step is omitted and what is called "convolution" is actually a "cross-correlation" between 2 matrices.



$$(A * B) * C = A * (B * C)$$

↑ the filter's mirroring, which we don't use in ML is used to have this property, useful in signal processing but not necessarily in ML.



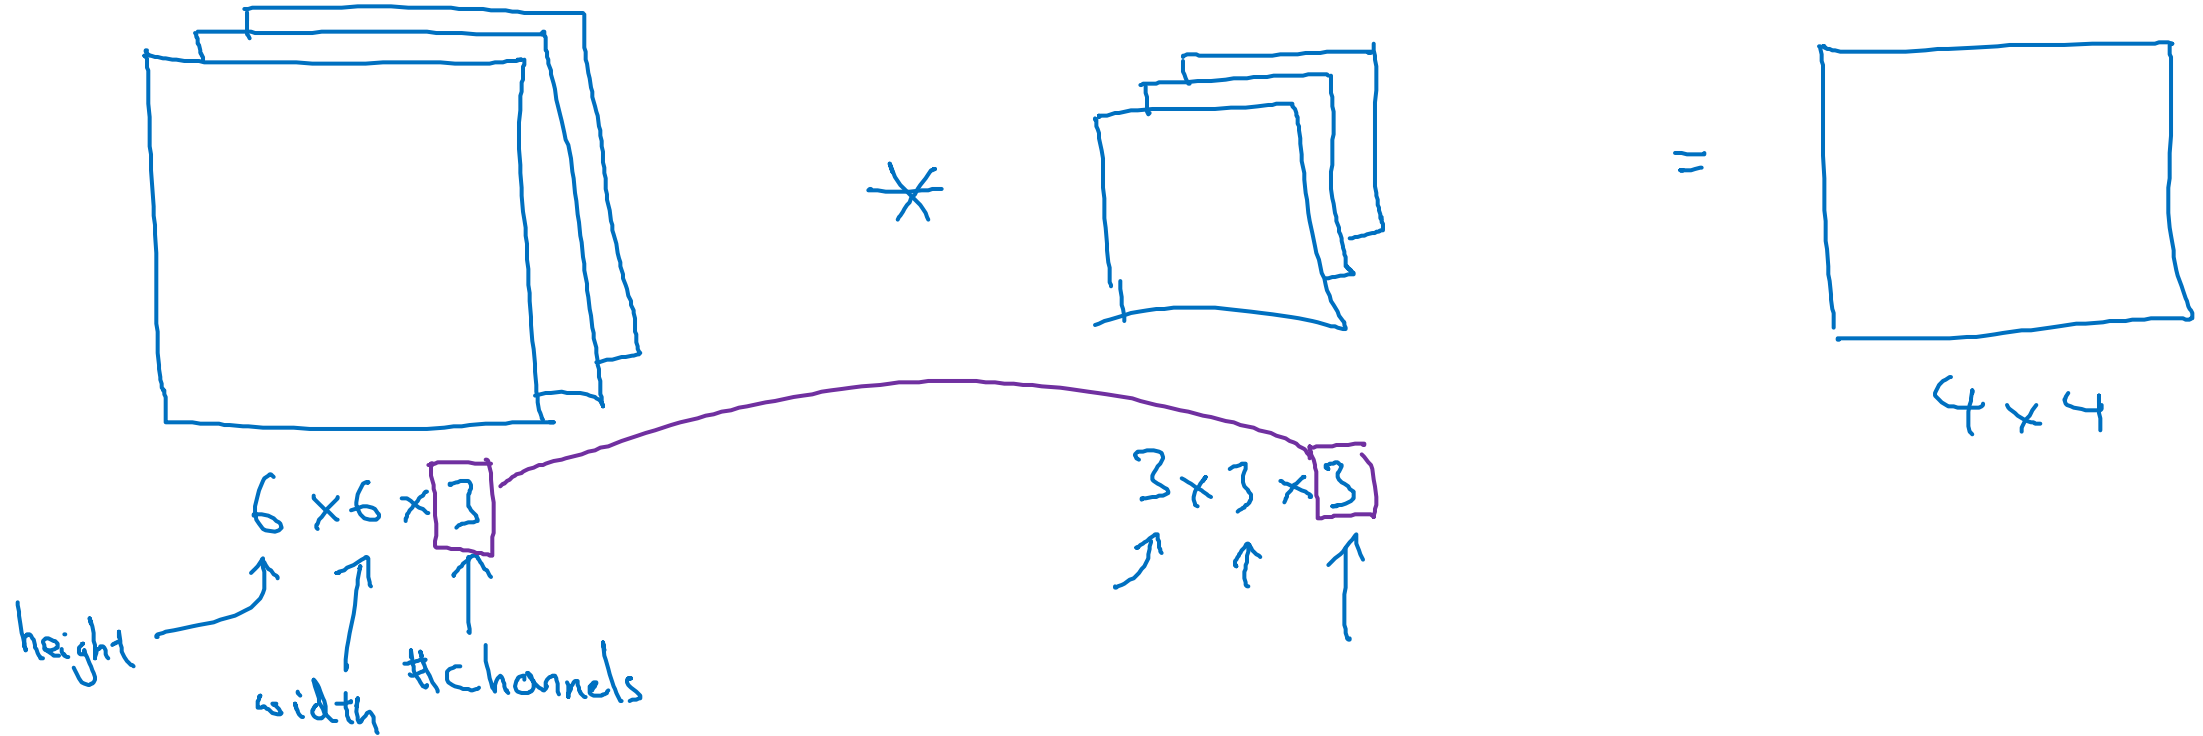
deeplearning.ai

# Convolutional Neural Networks


---

## Convolutions over volumes


# Convolutions on RGB images




# Convolutions on RGB image




$6 \times 6 \times 3$



\*






=





4 x 4

example 1:  
"filter that detects red  
edges in the R channel"



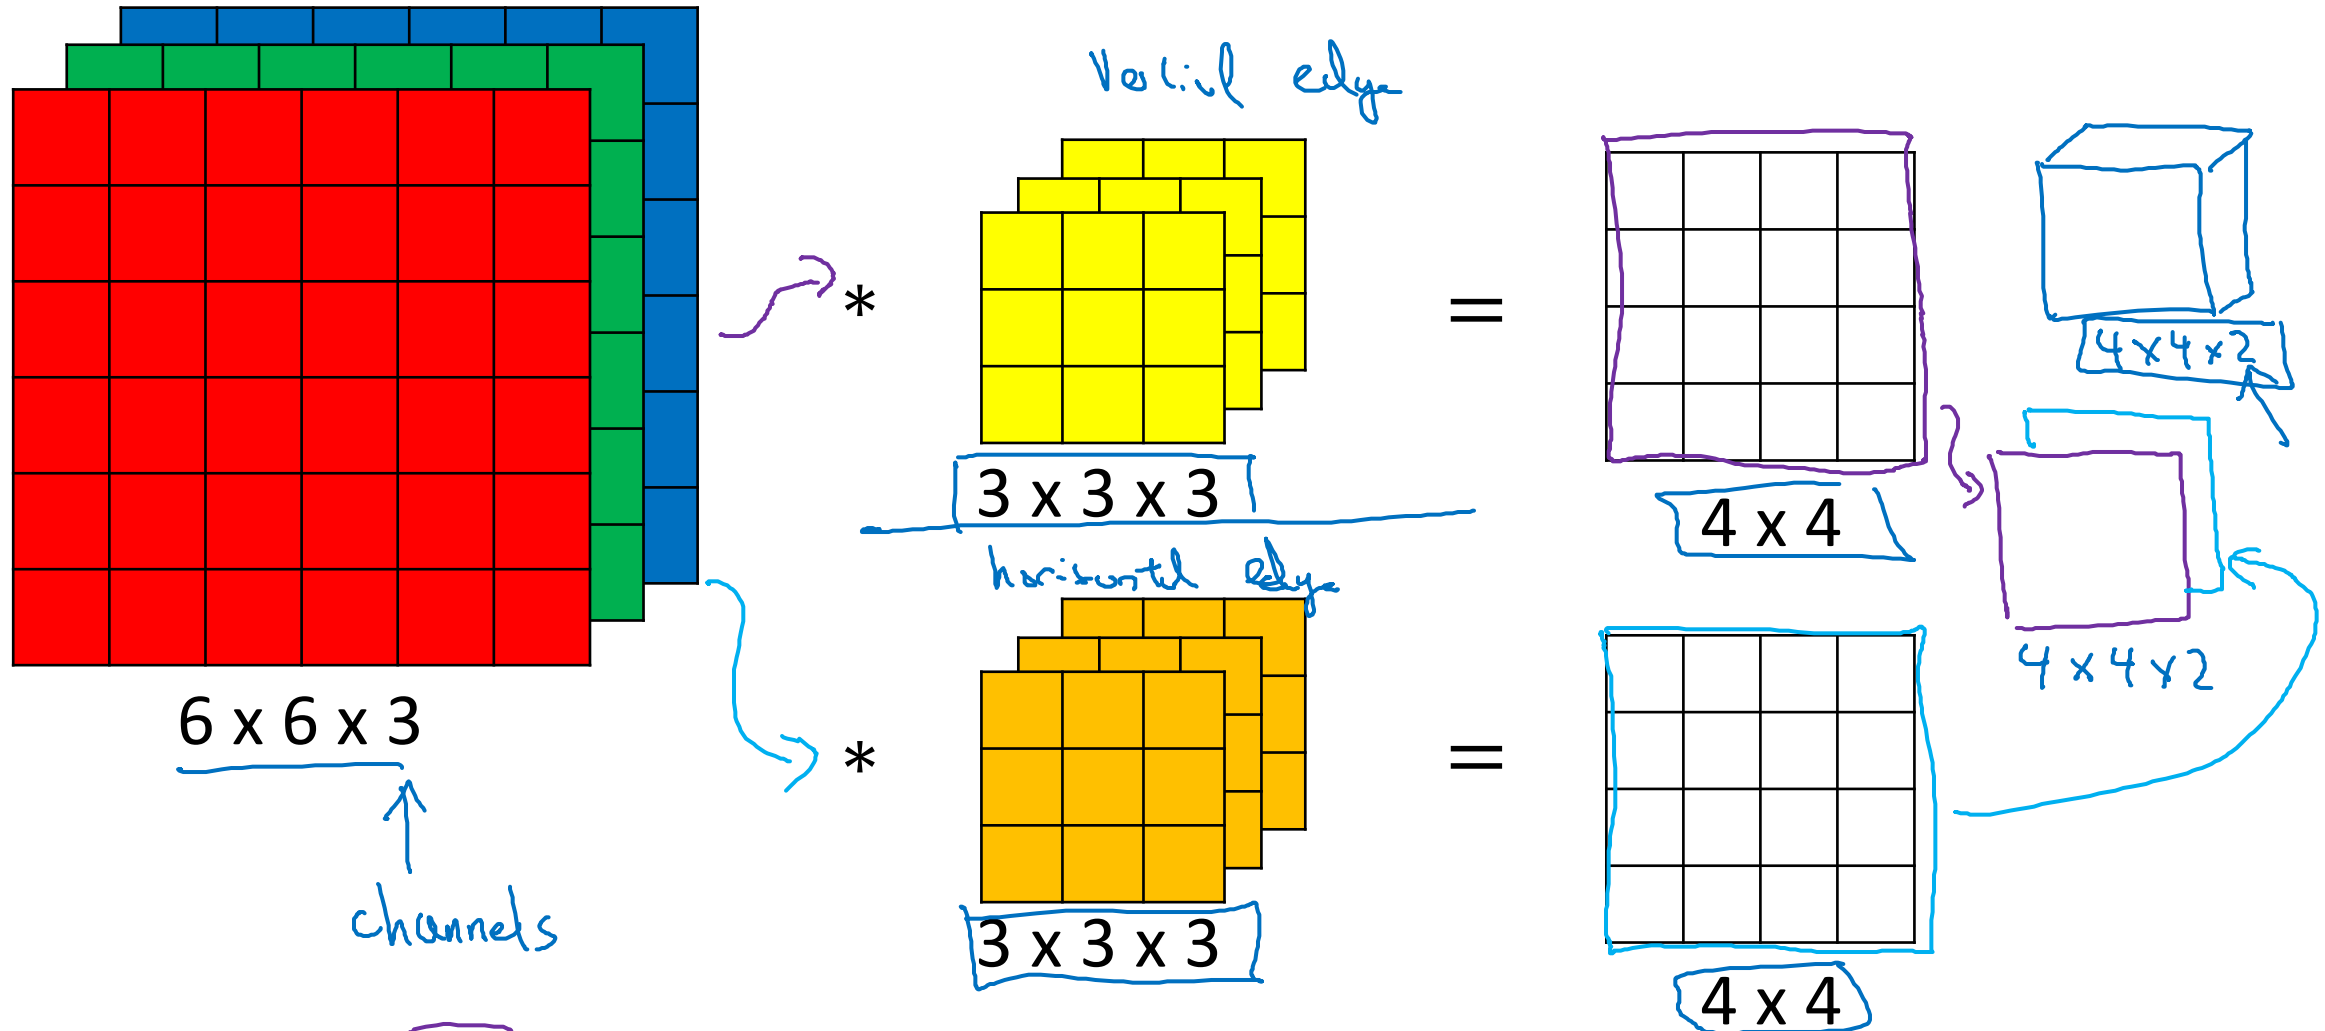
$\rightarrow 3 \times 3 \times 3$

example 2:  
"filter that detects  
vertical edges in  
any channel"



$\rightarrow 3 \times 3 \times 3$

# Multiple filters



Summary:  $n \times n \times n_c$   $\times$   $f \times f \times n_c$   $\rightarrow$   $\frac{n-f+1}{4} \times \frac{n-f+1}{4} \times n_c'$

$6 \times 6 \times 3$   $3 \times 3 \times 3$   $4 \times 4 \times 2$   $\uparrow$  #filters






deeplearning.ai

# Convolutional Neural Networks

---

One layer of a  
convolutional  
network

# Example of a layer




6 x 6 x 3

$a^{[0]}$


$$z^{[1]} = W^{[1]} a^{[0]} + b^{[1]}$$

$$a^{[1]} = g(z^{[1]})$$

\*  
\*




3 x 3 x 3




3 x 3 x 3

" $W^{[1]}$ "



6 x 6 x 3




$a^{[1]}$

4 x 4 x 2

→ Relu


" $W^{[1]} a^{[0]}$ "



4 x 4

+  $b_1$

↓ R



4 x 4


+  $b_2$

↓ R




4 x 4

↓



4 x 4




4 x 4 x 2

←  $a^{[1]}$

4 x 4 x 10

# Number of parameters in one layer

If you have 10 filters that are  $3 \times 3 \times 3$  in one layer of a neural network, how many parameters does that layer have?



# Summary of notation

If layer l is a convolution layer:

$f^{[l]}$  = filter size

$p^{[l]}$  = padding

$s^{[l]}$  = stride

$n_c^{[l]}$  = number of filters

→ Each filter is:  $f^{[l]} \times f^{[l]} \times n_c^{[l-1]}$

Activations:  $a^{[l]} \rightarrow n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}$

Weights:  $f^{[l]} \times f^{[l]} \times n_c^{[l-1]} \times n_c^{[l]}$

bias:  $n_c^{[l]} - (1, 1, 1, n_c^{[l]})$  ← #filters in layer l.

Input:  $n_H^{[l-1]} \times n_W^{[l-1]} \times n_c^{[l-1]}$

Output:  $n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}$

$$n_H^{[l]} = \left\lfloor \frac{n_H^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1 \right\rfloor$$

$A^{[l]} \rightarrow \underbrace{m \times n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}}_{n_c^{[l]} \times n_H^{[l]} \times n_W^{[l]}}$

#examples



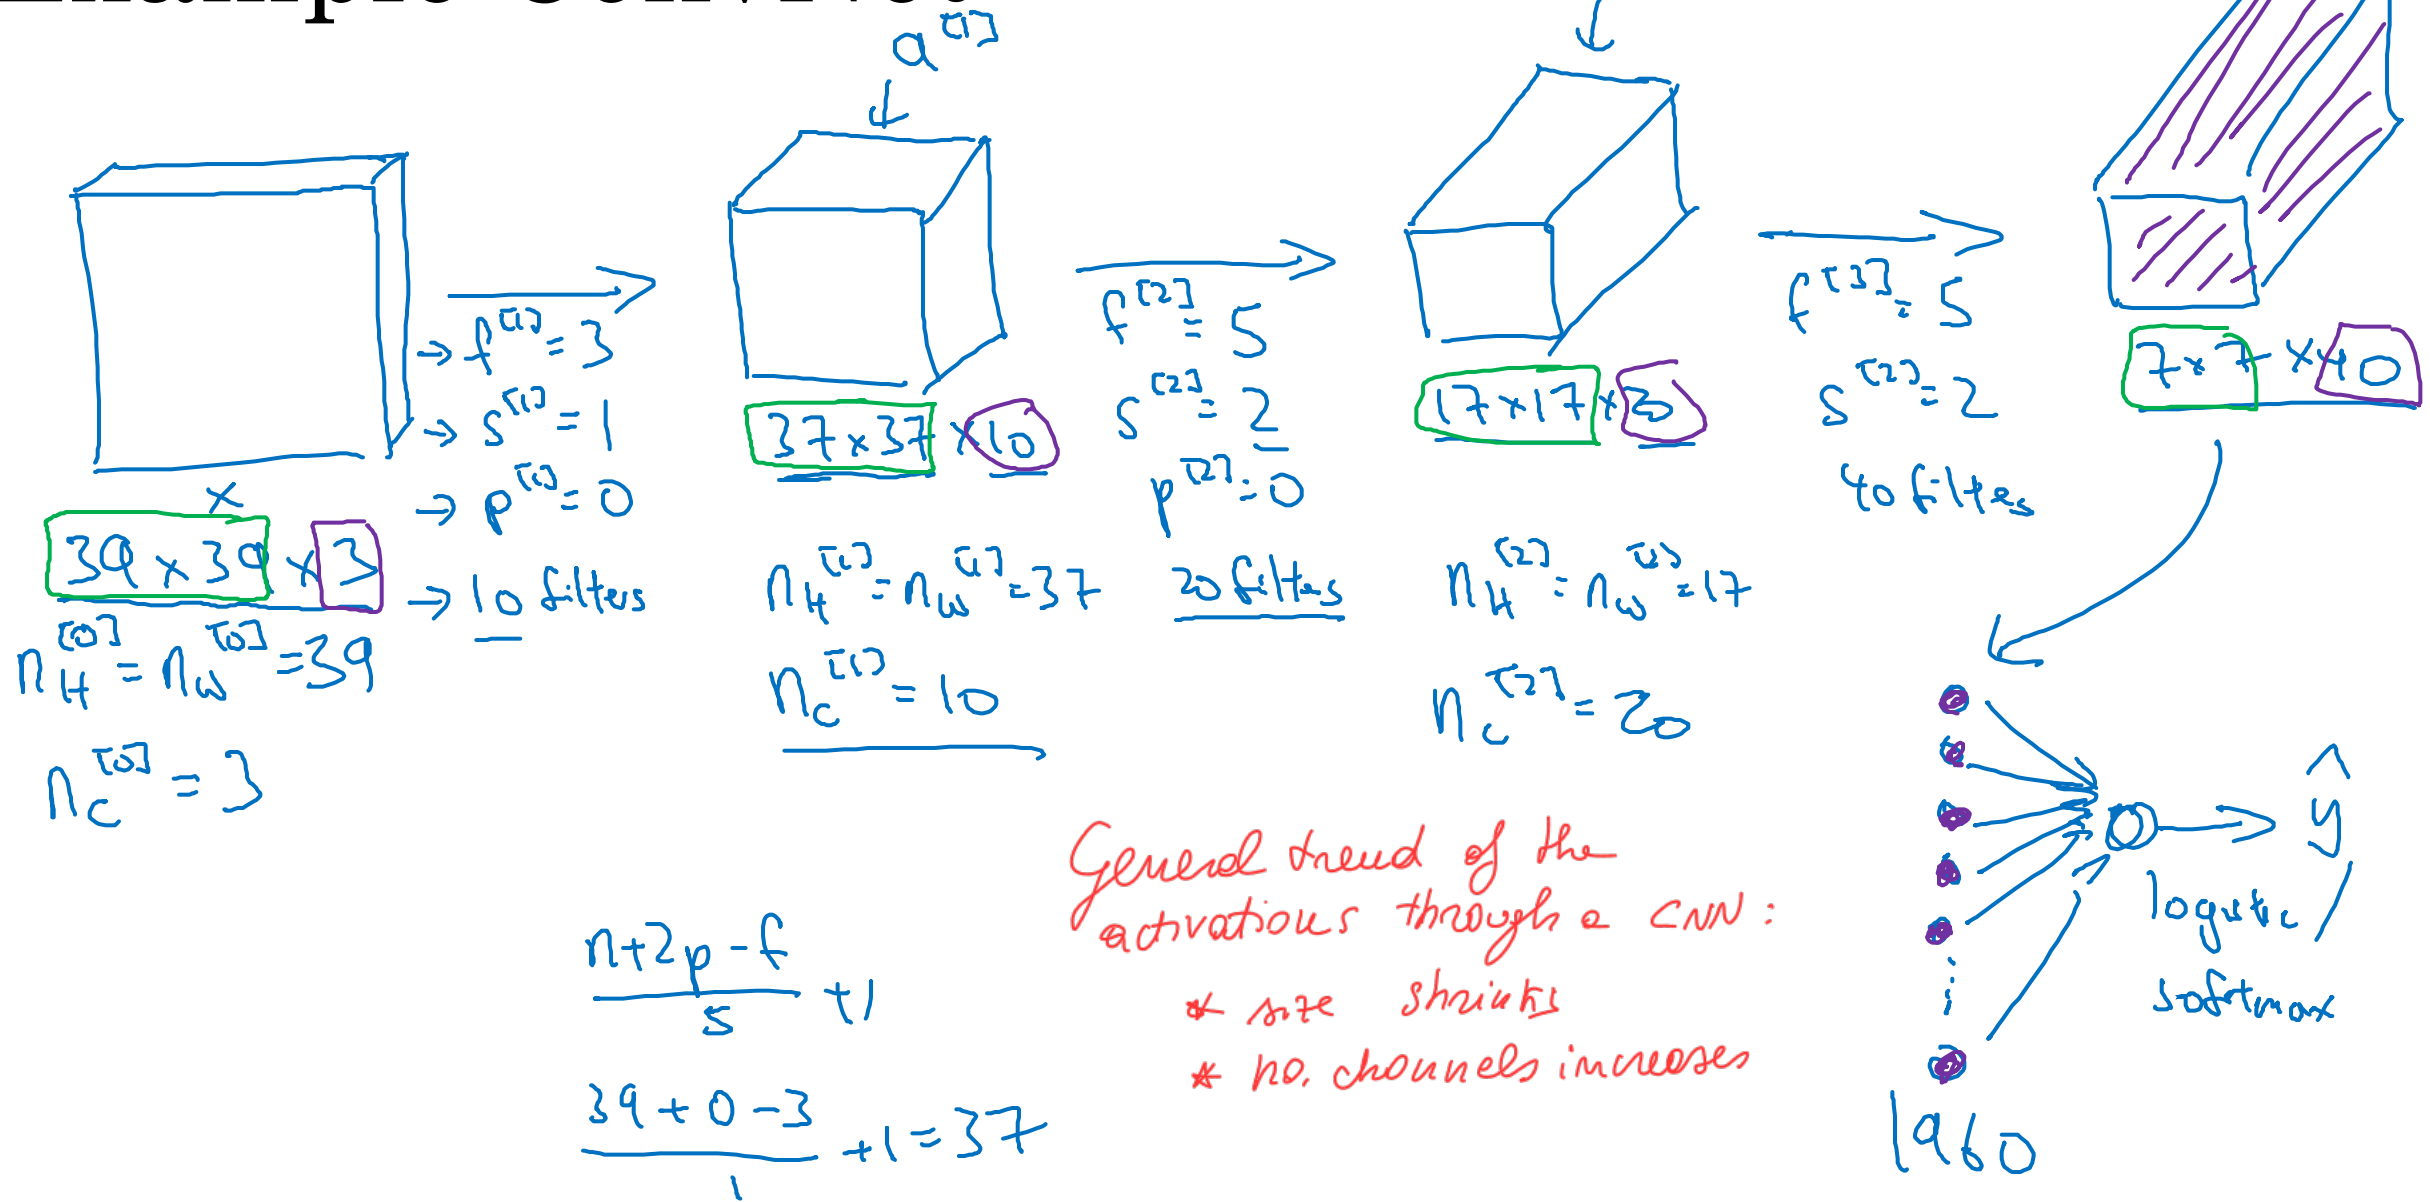
deeplearning.ai

# Convolutional Neural Networks

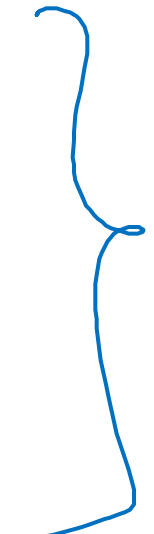
---

A simple convolution  
network example

# Example ConvNet



# Types of layer in a convolutional network:

- Convolution (conv) ←
  - Pooling (pool) ←
  - Fully connected (Fc) ←
- 





deeplearning.ai

# Convolutional Neural Networks

---

## Pooling layers

# Pooling layer: Max pooling

Shrinks activation maps by only retaining important info.

INTUITION: the activation of a filter indicates how much the detected feature appears in the input image.


1	3	2	1
2	9	1	1
1	3	2	3
5	6	1	2


By maxpooling the activation we discard some information while retaining

where in the image the feature is most prominent.

- performed on each channel independently
- no parameters to learn!

# Pooling layer: Max pooling



# Pooling layer: Average pooling

1	3	2	1
2	9	1	1
1	4	2	3
5	6	1	2



3.75	1.25
4	2

$$f=2$$

$$s=2$$

$$\underline{7 \times 7 \times 1000} \rightarrow 1 \times 1 \times 1000$$

# Summary of pooling

Hyperparameters:

f : filter size

$$f=2, s=2$$

s : stride

$$f=3, s=2$$

Max or average pooling

~~⇒ p: padding.~~

No parameters to learn!

$$n_H \times n_W \times \underline{n_C}$$



$$\left\lfloor \frac{n_H - f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n_W - f}{s} + 1 \right\rfloor \times \underline{n_C}$$



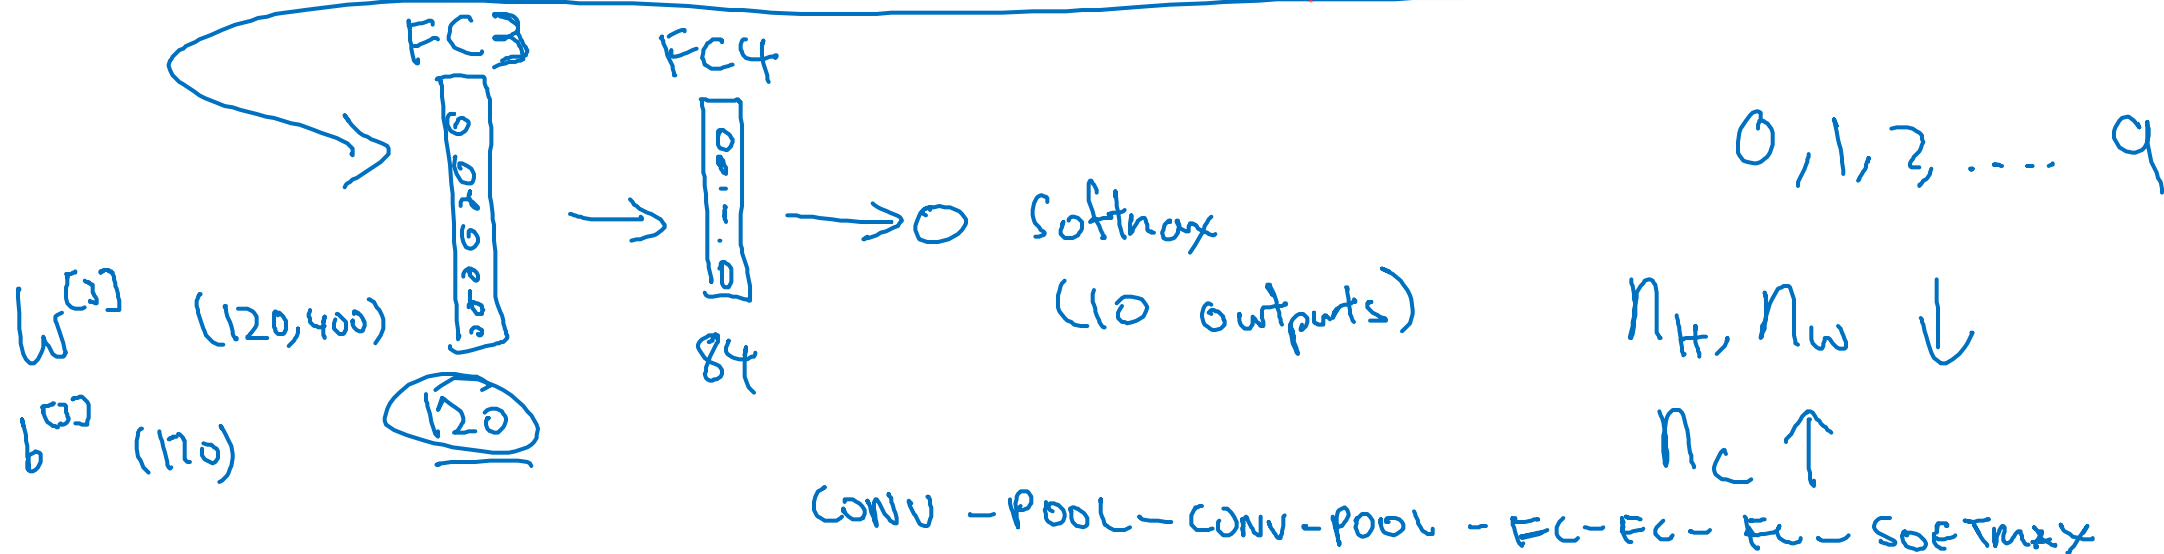
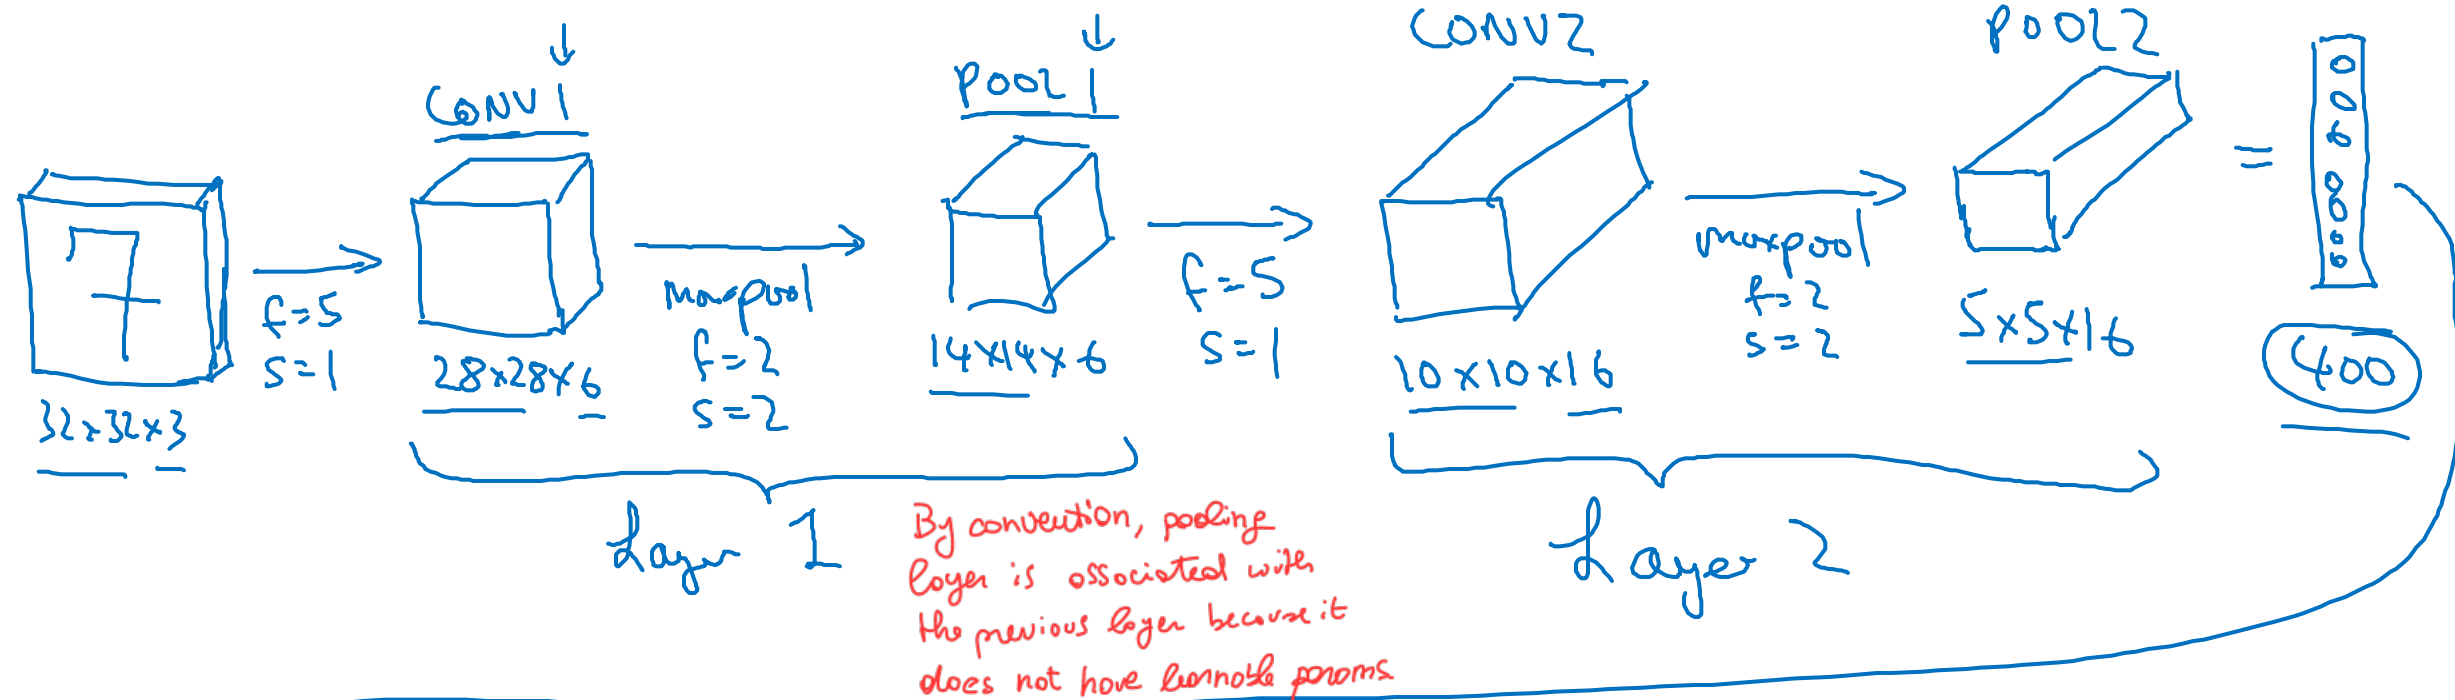
deeplearning.ai

# Convolutional Neural Networks

---

## Convolutional neural network example

# Neural network example (LeNet-5)





# Neural network example

	Activation shape	Activation Size	# parameters
Input:	(32,32,3)	3,072 $a^{[0]}$	0
CONV1 (f=5, s=1)	8 filters of size 5x5x3 (28,28,8)	<u>6,272</u>	608 ←
POOL1	(14,14,8)	<u>1,568</u>	0 ←
CONV2 (f=5, s=1)	(10,10,16)	<u>1,600</u>	3216 ←
POOL2	(5,5,16)	<u>400</u>	0 ←
FC3	(120,1)	<u>120</u>	48120 }
FC4	(84,1)	<u>84</u>	10164 }
Softmax	(10,1)	<u>10</u>	850





deeplearning.ai

# Convolutional Neural Networks

---

## Why convolutions?

# Why convolutions




$$5 \times 5 = 25$$

$$26$$

$$6 \times 26 = 156 \text{ parameters}$$

$$3,072 \times 4,704 \approx \underline{14M}$$

# Why convolutions



**Parameter sharing:** A feature detector (such as a vertical edge detector) that's useful in one part of the image is probably useful in another part of the image.


(As opposed to a FC layer connected to an image, which would release the same feature multiple times in different units).  
SPATIAL LOCALITY  $\Rightarrow$  also favors translation invariance. & reduces overfitting.

→ **Sparsity of connections:** In each layer, each output value depends only on a small number of inputs.



# Putting it together

Training set  $(x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})$ .



$$\text{Cost } J = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Use gradient descent to optimize parameters to reduce  $J$